

Homework 2

Tianyu Liu (301249861)

2023-10-23

Problem Set 7, Applications

1

(a)

```
AQ <- na.omit(airquality[,1:4])
AQ$TWcp = AQ$Temp*AQ$Wind
AQ$TWrat = AQ$Temp/AQ$Wind
head(AQ)
```

```
##   Ozone Solar.R Wind Temp  TWcp    TWrat
## 1    41     190   7.4   67 495.8 9.054054
## 2    36     118   8.0   72 576.0 9.000000
## 3    12     149  12.6   74 932.4 5.873016
## 4    18     313  11.5   62 713.0 5.391304
## 7    23     299   8.6   65 559.0 7.558140
## 8    19      99  13.8   59 814.2 4.275362
```

```
library(MASS)
```

```
#####
# Ridge regression using lm.ridge()
#
# Running sequence from 0 to 100 by a small increment may be a little expensive
# but probably covers range.
ridge <- lm.ridge(Ozone ~ ., lambda = seq(0, 100, .05), data = AQ)
select(ridge)
```

```
## modified HKB estimator is 0.3437551
## modified L-W estimator is 1.534077
## smallest value of GCV at 0.2
```

```
(coef.ri.best <- coef(ridge)[which.min(ridge$GCV), ])
```

```
##           Solar.R           Wind           Temp           TWcp
## -161.63123617    0.06284069    6.82529896    2.45651021   -0.10883212
##           TWrat
##      1.64685249
```

```
coef.ri.best
```

```
##           Solar.R           Wind           Temp           TWcp
## -161.63123617    0.06284069    6.82529896    2.45651021   -0.10883212
##           TWrat
##      1.64685249
```

(b)

```
# Compare MSPE to MSPE from lm
mod.lm <- lm(Ozone ~ ., data = AQ)
coef.ls <- coef(mod.lm)

length(which(coef.ls[2:6] - coef.ri.best[2:6] > 0)) / length(coef.ls[2:6])
```

```
## [1] 0.6
```

3 of the 5 parameter estimates (60%) are smaller than the least square estimate.

2

(a)

```
library(glmnet)
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 4.1-8
```

```

y<- AQ[,1]
x<- as.matrix(AQ[, c(2:6)])

# Fit LASSO by glmnet(y=, x=). Gaussian is default, but other families are available
# Function produces series of fits for many values of lambda.

lasso <- glmnet(y = y, x = x, family = "gaussian")

# cv.glmnet() uses crossvalidation to estimate optimal lambda

cv.lasso <- cv.glmnet(y = y, x = x, family = "gaussian")
cv.lasso

```

```

##
## Call:  cv.glmnet(x = x, y = y, family = "gaussian")
##
## Measure: Mean-Squared Error
##
##      Lambda Index Measure      SE Nonzero
## min  0.007    89   418.4 120.3        5
## 1se  8.651    12   529.2 137.4        2

```

(b)

```
coef(cv.lasso, s = cv.lasso$lambda.min)
```

```

## 6 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept) -183.90117408
## Solar.R      0.06354556
## Wind         8.88369468
## Temp         2.78797591
## TWcp         -0.13800612
## TWrat         1.43082832

```

```
coef(cv.lasso, s = cv.lasso$lambda.1se)
```

```

## 6 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept) -42.3947165
## Solar.R      .
## Wind         .
## Temp         0.8586877
## TWcp         .
## TWrat         1.8785311

```

Note that the 1se model shrunk 3 of the parameters (Solar and TWcp) and they weren't included in the model, while the λ -min model shrunk only 1 (Wind)

(c)

```
step <- step(
  object = lm(y ~ 1, data = AQ), scope = list(upper = mod.lm), direction = "both",
  k = log(nrow(AQ)), trace = 0)
coef.step <- coef(step)
coef.step
```

```
## (Intercept)      TWrat      Temp      Solar.R
## -93.30421016   2.86326061   1.25230788   0.05959571
```

The λ -min model has 4 of the 5 variables, the 1se model has 2 of the 5, and the hybrid stepwise has 3 of the 5.

3

(a),(b),(c),(d)

```
set.seed(2928893)
### Let's define a function for constructing CV folds
get.folds <- function(n, K) {
  ### Get the appropriate number of fold labels
  n.fold <- ceiling(n / K) # Number of observations per fold (rounded up)
  fold.ids.raw <- rep(1:K, times = n.fold) # Generate extra labels
  fold.ids <- fold.ids.raw[1:n] # Keep only the correct number of labels

  ### Shuffle the fold labels
  folds.rand <- fold.ids[sample.int(n)]

  return(folds.rand)
}

get.MSPE <- function(Y, Y.hat) {
  return(mean((Y - Y.hat)^2))
}

### Number of folds
K <- 10

### Construct folds
n <- nrow(AQ) # Sample size
folds <- get.folds(n, K)

### Create a container for MSPEs. Let's include ordinary least-squares
### regression for reference
all.models <- c("Ridge", "LASSO-Min", "LASSO-1se")
all.MSPEs <- array(0, dim = c(K, length(all.models)))
colnames(all.MSPEs) <- all.models

### Begin cross-validation
for (i in 1:K) {
  ### Split data
  data.train <- AQ[folds != i, ]
  data.valid <- AQ[folds == i, ]
  n.train <- nrow(data.train)

  ### Get response vectors
  Y.train <- data.train$Ozone
  Y.valid <- data.valid$Ozone

  #####
  ### Let's do ridge regression. This model is fit using the ###
  ### lm.ridge() function in the MASS package. We will need to make a ###
  ### list of candidate lambda values for the function to choose ###
  ### from. Prediction also has some extra steps, but we'll discuss ###
}
```

```

### that when we get there.                                     ###
#####

### Make a list of lambda values. The lm.ridge() function will
### then choose the best value from this list. Use the seq()
### function to create an equally-spaced list.
lambda.vals <- seq(from = 0, to = 100, by = 0.05)

### Use the lm.ridge() function to fit a ridge regression model. The
### syntax is almost identical to the lm() function, we just need
### to set lambda equal to our list of candidate values.
fit.ridge <- lm.ridge(Ozone ~ .,
  lambda = lambda.vals,
  data = data.train
)

### To get predictions, we need to evaluate the fitted regression
### equation directly (sadly, no predict() function to do this for us).
### You could do this using a for loop if you prefer, but there is
### a shortcut which uses matrix-vector multiplication. The syntax
### for this multiplication method is much shorter.

### Get best lambda value and its index
### Note: Best is chosen according to smallest GCV value. We can
### get GCV from a ridge regression object using $GCV
ind.min.GCV <- which.min(fit.ridge$GCV)
lambda.min <- lambda.vals[ind.min.GCV]

### Get coefficients corresponding to best lambda value
### We can get the coefficients for every value of lambda using
### the coef() function on a ridge regression object
all.coefs.ridge <- coef(fit.ridge)
coef.min <- all.coefs.ridge[ind.min.GCV, ]

### We will multiply the dataset by this coefficients vector, but
### we need to add a column to our dataset for the intercept and
### create indicators for our categorical predictors. A simple
### way to do this is using the model.matrix() function from last
### week.
matrix.valid.ridge <- model.matrix(Ozone ~ ., data = data.valid)

### Now we can multiply the data by our coefficient vector. The
### syntax in R for matrix-vector multiplication is %*%. Note that,
### for this type of multiplication, order matters. That is,
###  $A \%*\% B \neq B \%*\% A$ . Make sure you do data %*% coefficients.
### For more information, see me in a Q&A session or, better still,
### take a course on linear algebra (it's really neat stuff)
pred.ridge <- matrix.valid.ridge %*% coef.min

### Now we just need to calculate the MSPE and store it
MSPE.ridge <- get.MSPE(Y.valid, pred.ridge)
all.MSPEs[i, "Ridge"] <- MSPE.ridge

```

```
#####
### Now we can do the LASSO. This model is fit using the glmnet() ###
### or cv.glmnet() functions in the glmnet package. LASSO also has ###
### a tuning parameter, Lambda, which we have to choose. ###
### Fortunately, the cv.glmnet() function does CV internally, and ###
### lets us automatically find the 'best' value of lambda. ###
#####
```

```
### The cv.glmnet() function has different syntax from what we're
### used to. Here, we have to provide a matrix with all of our
### predictors, and a vector of our response. LASSO handles
### the intercept differently, so we want to make sure our data
### matrix does not include an intercept (then let cv.glmnet() add
### an intercept later). Unfortunately, the model.matrix() function
### gets confused if we ask it to construct indicators for our
### categorical predictors without also including an intercept.
### A simple way to fix this is to create the data matrix with an
### intercept, then delete the intercept.
matrix.train.raw <- model.matrix(Ozone ~ ., data = data.train)
matrix.train <- matrix.train.raw[, -1]
```

```
### The cv.glmnet() function creates a list of Lambda values, then
### does CV internally to choose the 'best' one. 'Best' can refer to
### either the value of lambda which gives the smallest CV-MSPE
### (called the min rule), or the value of lambda which gives the
### simplest model that gives CV-MSPE close to the minimum (called
### the 1se rule). The cv.glmnet() function gets both of these
### lambda values.
all.LASSOs <- cv.glmnet(x = matrix.train, y = Y.train)
```

```
### Get both 'best' lambda values using $lambda.min and $lambda.1se
lambda.min <- all.LASSOs$lambda.min
lambda.1se <- all.LASSOs$lambda.1se
```

```
### cv.glmnet() has a predict() function (yay!). This predict function
### also does other things, like get the coefficients, or tell us
### which predictors get non-zero coefficients. We are also able
### to specify the value of lambda for which we want our output
### (remember that, with ridge, we got a matrix of coefficients and
### had to choose the row matching our lambda). Strangely, the name
### of the input where we specify our value of lambda is s.
```

```
### Get the coefficients for our two 'best' LASSO models
coef.LASSO.min <- predict(all.LASSOs, s = lambda.min, type = "coef")
coef.LASSO.1se <- predict(all.LASSOs, s = lambda.1se, type = "coef")
```

```
### Get which predictors are included in our models (i.e. which
### predictors have non-zero coefficients)
included.LASSO.min <- predict(all.LASSOs,
  s = lambda.min,
```

```

    type = "nonzero"
  )
included.LASSO.1se <- predict(all.LASSOs,
  s = lambda.1se,
  type = "nonzero"
)

### Get predictions from both models on the validation fold. First,
### we need to create a predictor matrix from the validation set.
### Remember to include the intercept in model.matrix(), then delete
### it in the next step.
matrix.valid.LASSO.raw <- model.matrix(Ozone ~ ., data = data.valid)
matrix.valid.LASSO <- matrix.valid.LASSO.raw[, -1]
pred.LASSO.min <- predict(all.LASSOs,
  newx = matrix.valid.LASSO,
  s = lambda.min, type = "response"
)
pred.LASSO.1se <- predict(all.LASSOs,
  newx = matrix.valid.LASSO,
  s = lambda.1se, type = "response"
)

### Calculate MSPEs and store them
MSPE.LASSO.min <- get.MSPE(Y.valid, pred.LASSO.min)
all.MSPEs[i, "LASSO-Min"] <- MSPE.LASSO.min

MSPE.LASSO.1se <- get.MSPE(Y.valid, pred.LASSO.1se)
all.MSPEs[i, "LASSO-1se"] <- MSPE.LASSO.1se
}

all.MSPEs

```

```

##           Ridge LASSO-Min LASSO-1se
## [1,] 280.6840 314.2722 415.7364
## [2,] 363.7170 316.8250 236.9153
## [3,] 539.3697 586.1201 717.2987
## [4,] 108.3192 114.6189 157.0742
## [5,] 185.4001 190.1739 323.5662
## [6,] 371.7604 373.4553 668.8504
## [7,] 219.6295 217.3012 366.8750
## [8,] 587.6590 585.7460 643.3282
## [9,] 1042.9887 1045.5646 1125.8526
## [10,] 725.3818 650.6659 377.7573

```

```

avg.MSPEs = colMeans(all.MSPEs)
avg.MSPEs

```

```

##           Ridge LASSO-Min LASSO-1se
## 442.4909 439.4743 503.3254

```


On average, ridge regression produces the smallest prediction error.

(e)

```
library(stringr)

all.MSPEs.LS.step <- array(0, dim = c(K, 2))
colnames(all.MSPEs.LS.step) <- c("LS", "hybrid stepwise")

coefs.sw <- matrix(NA, nrow = K, ncol = 11)

for (i in 1:K) {
  ### Split data
  data.train <- AQ[folds != i, ]
  data.valid <- AQ[folds == i, ]
  n.train <- nrow(data.train)

  ### Get response vectors
  Y.train <- data.train$Ozone
  Y.valid <- data.valid$Ozone

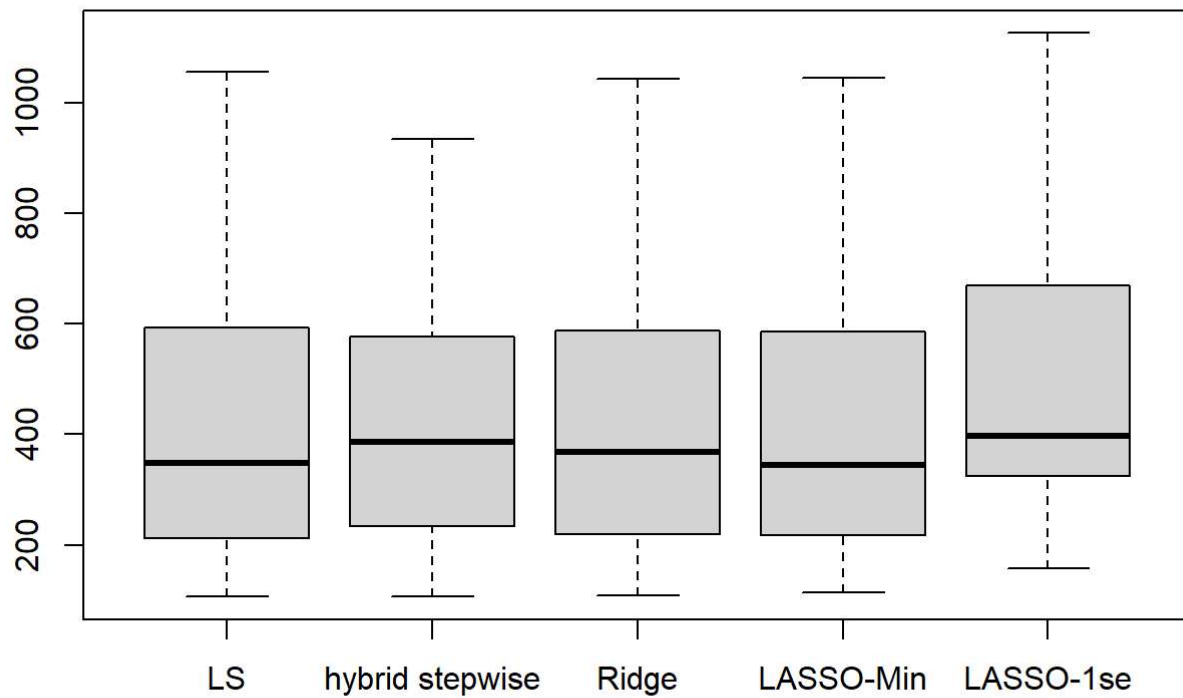
  #####
  ### First, let's quickly do LS so we have a reference point for ###
  ### how well the other models do. #####
  #####

  fit.ls <- lm(Ozone ~ ., data = data.train)
  pred.ls <- predict(fit.ls, newdata = data.valid)
  MSPE.ls <- get.MSPE(Y.valid, pred.ls)
  all.MSPEs.LS.step[i, "LS"] <- MSPE.ls

  step <- step(
    object = lm(Ozone ~ 1, data = data.train), scope = list(upper = fit.ls), direction = "both",
    k = log(nrow(data.train)), trace = 0)
  pred.sw <- predict(step, as.data.frame(data.valid))
  MSPE.sw <- get.MSPE(Y.valid, pred.sw)
  all.MSPEs.LS.step[i, "hybrid stepwise"] <- MSPE.sw
}

all.MSPEs.all = cbind(all.MSPEs.LS.step, all.MSPEs)
boxplot(all.MSPEs.all, main = paste0("CV MSPEs over ", K, " folds"))
```

CV MSPEs over 10 folds

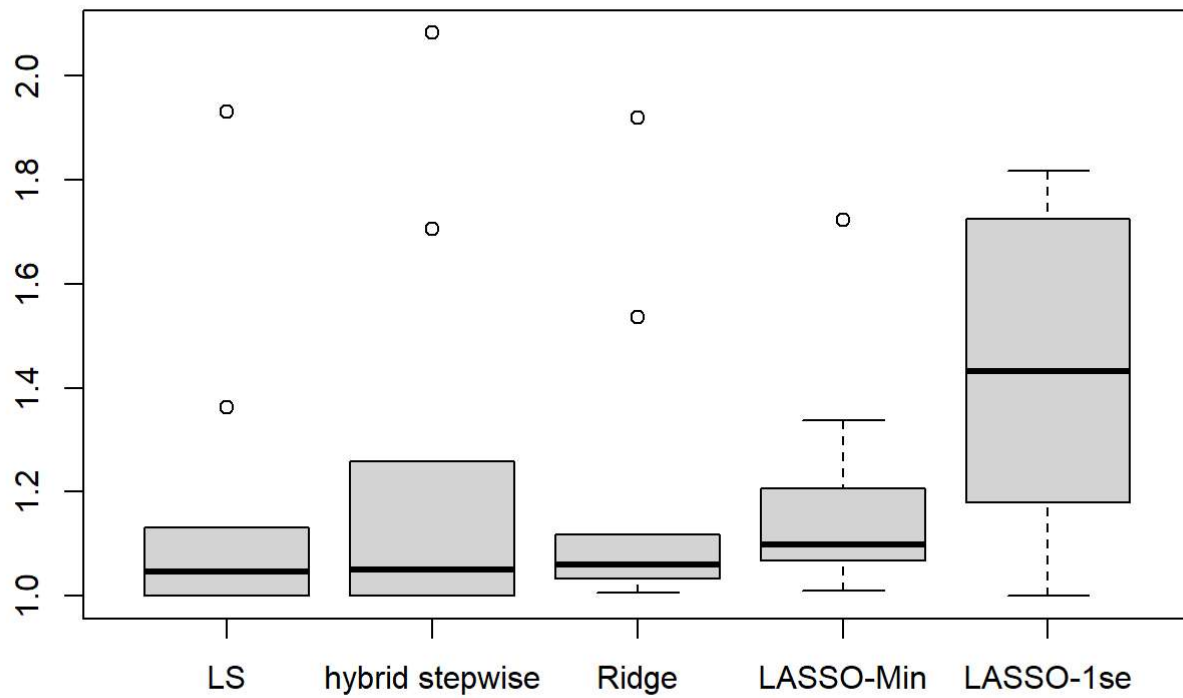


There isn't any big differences. The least-square and ridge models have lower MSPE but higher variance, whereas the LASSO-1se model has the lowest variance (but has a clear outlier).

(f)

```
all.RMSPEs <- apply(all.MSPEs.all, 1, function(W) {  
  best <- min(W)  
  return(W / best)  
})  
all.RMSPEs <- t(all.RMSPEs)  
boxplot(all.RMSPEs, main = paste0("CV RMSPEs over ", K, " folds"))
```

CV RMSPEs over 10 folds



The

least-square and ridge models behave the best still.

Problem Set 8, Applications (Ozone Data)

1

(a),(b),(c)

```
library(pls)
```

```
##  
## Attaching package: 'pls'
```

```
## The following object is masked from 'package:stats':  
##  
## loadings
```

```

all.MSPEs.pls <- array(0, dim = c(K, 1))
colnames(all.MSPEs.pls) <- c("PLS")

n_comps = array(0, dim = c(K,1))

for (i in 1:K) {
  ### Split data
  data.train <- AQ[folds != i, ]
  data.valid <- AQ[folds == i, ]
  n.train <- nrow(data.train)

  ### Get response vectors
  Y.train <- data.train$Ozone
  Y.valid <- data.valid$Ozone

  ### Now, Let's do PLS using the plsr() function. The syntax is
  ### very similar to lm(). If we set validation = "CV", the plsr()
  ### function will do its own internal CV, and give MSPEs for each
  ### number of components. We can then use this to choose how many
  ### componenets to keep when doing prediction on the validation
  ### fold. We can use an optional input called segments to specify
  ### how many folds we want plsr() to use for its internal CV
  ### (default is 10).
  fit_pls <- plsr(Ozone ~ .,
    data = data.train, validation = "CV",
    segments = 5
  )

  ### Investigate the fitted PLS model. Comment out the next two
  ### lines when running a CV loop

  ### The summary function gives us lots of information about how
  ### errors change as we increase the number of components
  # summary(fit.pls)

  ### The validationplot() function shows how MSPE from the internal
  ### CV of plsr() changes with the number of included components.
  # validationplot(fit.pls)

  ### Get the best model from PLS. To do this, we need to find the model
  ### that minimizes MSPE for the plsr() function's internal CV. It
  ### takes a few steps, but all the information we need is contained
  ### in the output of plsr().
  CV_pls <- fit_pls$validation
  pls_comps <- CV_pls$PRESS
  n_comps[i] <- which.min(pls_comps)

  ### Get predictions and calculate MSPE on the validation fold
  ### Set ncomps equal to the optimal number of components
  pred.pls <- predict(fit_pls, data.valid, ncomp = n_comps[i])
  MSPE.pls <- get.MSPE(Y.valid, pred.pls)
  all.MSPEs.pls[i, "PLS"] <- MSPE.pls
}

```

```
}  
  
n_comps
```

```
##      [,1]  
## [1,]    5  
## [2,]    3  
## [3,]    3  
## [4,]    3  
## [5,]    5  
## [6,]    4  
## [7,]    3  
## [8,]    3  
## [9,]    5  
## [10,]   3
```

```
all.MSPEs.pls
```

```
##      PLS  
## [1,] 260.4455  
## [2,] 139.8033  
## [3,] 513.6342  
## [4,] 137.6674  
## [5,] 192.1326  
## [6,] 364.8359  
## [7,] 265.6754  
## [8,] 655.6349  
## [9,] 1055.8045  
## [10,] 699.0148
```

```
mean(all.MSPEs.pls)
```

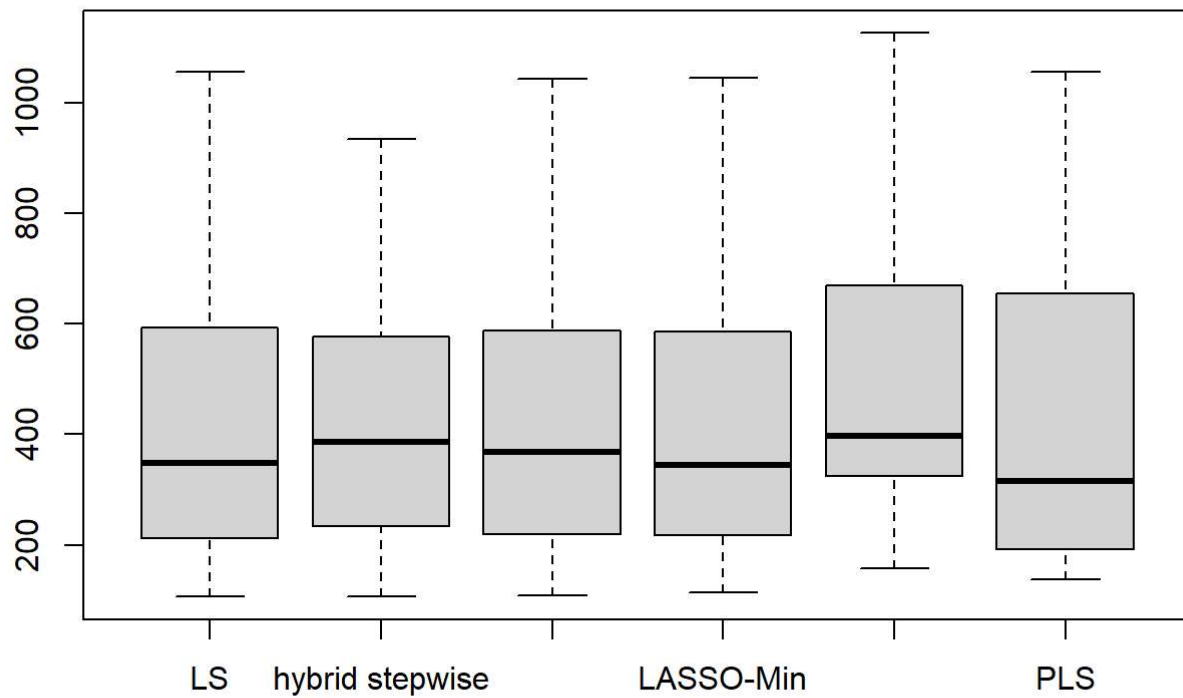
```
## [1] 428.4649
```

The optimal number of components for each of the 10 folds range from 3 to 5 (5 just means that each variable is its own component). The average prediction error is 433.

(d)

```
boxplot(cbind(all.MSPEs.all, all.MSPEs.pls), main = paste0("CV MSPEs over ", K, " folds"))
```

CV MSPEs over 10 folds




Compared with other models, the partial least squares estimate has the lowest MSPE and variance, but seems to be heavily skewed.

(e)

```
all.RMSPEs.all <- apply(cbind(all.MSPEs.all, all.MSPEs.pls), 1, function(W) {  
  best <- min(W)  
  return(W / best)  
})  
all.RMSPEs.all <- t(all.RMSPEs.all)  
boxplot(all.RMSPEs.all, main = paste0("CV RMSPEs over ", K, " folds"))
```

CV RMSPEs over 10 folds



The

PLS prediction outperforms every other model, according to the relative prediction error. The LS and ridge models are also comparable to PLS.

Problem Set 9, Concepts

1

(a)

β_o is the intercept value when the explanatory variable, X , is in region 0 of the step function (region used as the baseline) in the regression model.

(b)

β_k is the difference of intercept values when X is in the last region (region k) and β_o , the baseline region.

Problem Set 10, Applications

(a)

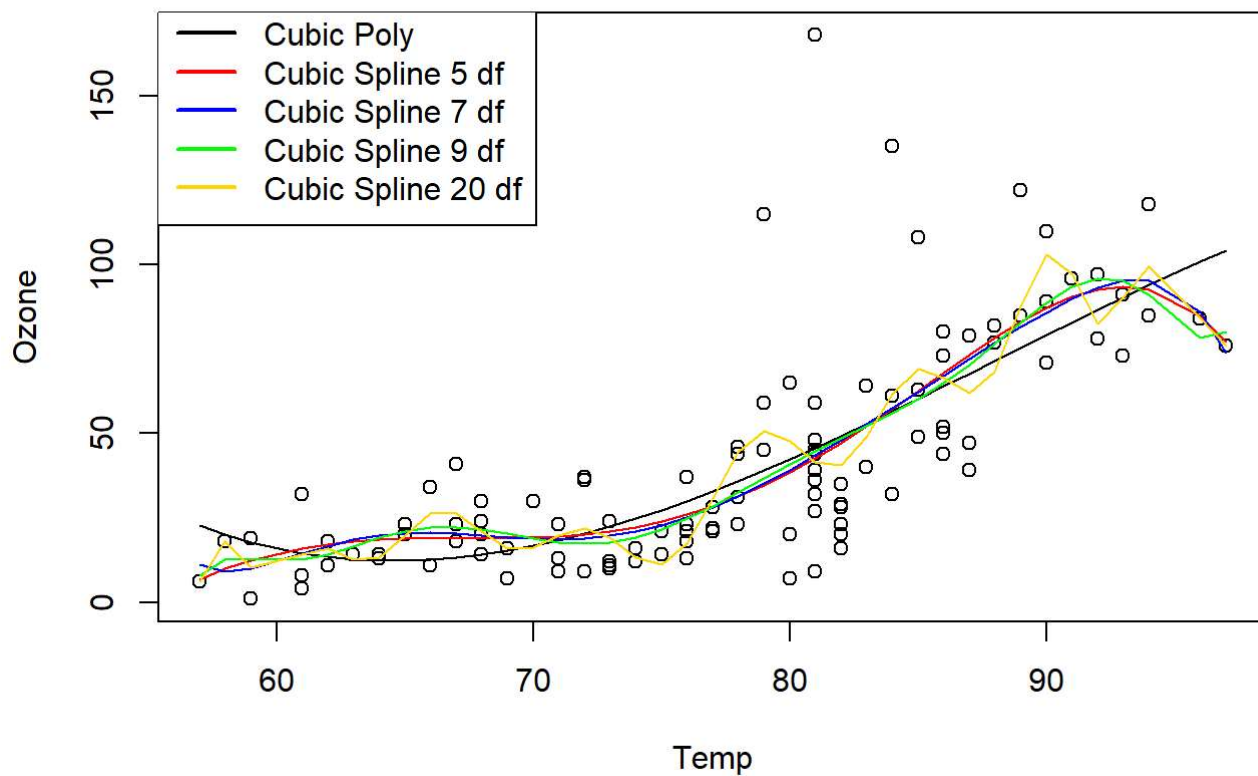
```
### Fit polynomial regression models
fit.poly.3 <- lm(Ozone ~ poly(Temp, degree = 3), data = AQ)

### Fit basis splines
library(splines)

fit.basis.5 <- lm(Ozone ~ bs(Temp, degree = 5), data = AQ)
fit.basis.7 <- lm(Ozone ~ bs(Temp, degree = 7), data = AQ)
fit.basis.9 <- lm(Ozone ~ bs(Temp, degree = 9), data = AQ)
fit.basis.20 <- lm(Ozone ~ bs(Temp, degree = 20), data = AQ)

### Predicting
Temp.sort <- data.frame(Temp = sort(AQ$Temp))
pred.poly.3 <- predict(fit.poly.3, Temp.sort)
pred.basis.5 <- predict(fit.basis.5, Temp.sort)
pred.basis.7 <- predict(fit.basis.7, Temp.sort)
pred.basis.9 <- predict(fit.basis.9, Temp.sort)
pred.basis.20 <- predict(fit.basis.20, Temp.sort)

### Plots
with(AQ, plot(Temp, Ozone))
lines(Temp.sort$Temp, pred.poly.3)
lines(Temp.sort$Temp, pred.basis.5, col = 'red')
lines(Temp.sort$Temp, pred.basis.7, col = 'blue')
lines(Temp.sort$Temp, pred.basis.9, col = 'green')
lines(Temp.sort$Temp, pred.basis.20, col = 'gold')
legend(x = 55, y = 180, legend = c(
  "Cubic Poly", "Cubic Spline 5 df",
  "Cubic Spline 7 df", "Cubic Spline 9 df", "Cubic Spline 20 df"
),
lty = "solid", col = c('black', 'red', 'blue', 'green', 'gold'), lwd = 2
)
```

(b)

Looking at the graph, the cubic polynomial seems to have the most bias.

(c)

The 9-df and 20-df splines seem to overfit, because there is some fluctuation towards the left end of the graph where Temp is low. Also for the 20-df spline there is too much fluctuation in the middle.

(d)

I would choose the 5-df cubic spline. It seems to capture the data relatively well with no clear overfitting, and no random end-point fluctuations.